
You are currently looking at **version 1.2** of this notebook. To download notebooks and datafiles, as well as get help on Jupyter notebooks in the Coursera platform, visit the [Jupyter Notebook FAQ](https://www.coursera.org/learn/python-data-analysis/resources/0dhYG) (<https://www.coursera.org/learn/python-data-analysis/resources/0dhYG>) course resource.

Assignment 2 - Pandas Introduction

All questions are weighted the same in this assignment.

Part 1

The following code loads the olympics dataset (olympics.csv), which was derived from the Wikipedia entry on [All Time Olympic Games Medals](https://en.wikipedia.org/wiki/All-time_Olympic_Games_medal_table) (https://en.wikipedia.org/wiki/All-time_Olympic_Games_medal_table), and does some basic data cleaning.

The columns are organized as # of Summer games, Summer medals, # of Winter games, Winter medals, total # number of games, total # of medals. Use this dataset to answer the questions below.

```
In [1]: import pandas as pd

df = pd.read_csv('olympics.csv', index_col=0, skiprows=1)

for col in df.columns:
    if col[:2]=='01':
        df.rename(columns={col:'Gold'+col[4:]}, inplace=True)
    if col[:2]=='02':
        df.rename(columns={col:'Silver'+col[4:]}, inplace=True)
    if col[:2]=='03':
        df.rename(columns={col:'Bronze'+col[4:]}, inplace=True)
    if col[:1]=='#':
        df.rename(columns={col:'#'+col[1:]}, inplace=True)

names_ids = df.index.str.split('\s\(') # split the index by '('

df.index = names_ids.str[0] # the [0] element is the country name (new index)
df['ID'] = names_ids.str[1].str[:3] # the [1] element is the abbreviation or ID (take first 3 characters from that)

df = df.drop('Totals')
df.head()
```

Out[1]:

	# Summer	Gold	Silver	Bronze	Total	# Winter	Gold.1	Silver.1	Bronze.1	Tota
Afghanistan	13	0	0	2	2	0	0	0	0	0
Algeria	12	5	2	8	15	3	0	0	0	0
Argentina	23	18	24	28	70	18	0	0	0	0
Armenia	5	1	2	9	12	6	0	0	0	0
Australasia	2	3	4	5	12	0	0	0	0	0

Question 0 (Example)

What is the first country in df?

This function should return a Series.

```
In [2]: # You should write your whole answer within the function provided. The a
        # utograder will call
        # this function and compare the return value against the correct solutio
        # n value
        def answer_zero():
            # This function returns the row for Afghanistan, which is a Series o
            # bject. The assignment
            # question description will tell you the general format the autograd
            # er is expecting
            return df.iloc[0]

        # You can examine what your function returns by calling it in the cell.
        # If you have questions
        # about the assignment formats, check out the discussion forums for any
        # FAQs
        answer_zero()
```

```
Out[2]: # Summer          13
        Gold              0
        Silver            0
        Bronze            2
        Total             2
        # Winter          0
        Gold.1            0
        Silver.1          0
        Bronze.1          0
        Total.1           0
        # Games          13
        Gold.2            0
        Silver.2          0
        Bronze.2          2
        Combined total    2
        ID                AFG
        Name: Afghanistan, dtype: object
```

Question 1

Which country has won the most gold medals in summer games?

This function should return a single string value.

```
In [3]: def answer_one():
        return df['Gold'].idxmax()
        answer_one()
```

```
Out[3]: 'United States'
```

Question 2

Which country had the biggest difference between their summer and winter gold medal counts?

This function should return a single string value.

```
In [4]: def answer_two():
        df_copy = df
        df_copy['diff'] = (df_copy['Gold'] - df_copy['Gold.1']).abs()
        return df_copy['diff'].idxmax()

        answer_two()
```

```
Out[4]: 'United States'
```

Question 3

Which country has the biggest difference between their summer gold medal counts and winter gold medal counts relative to their total gold medal count?

$$\frac{\text{Summer Gold} - \text{Winter Gold}}{\text{Total Gold}}$$

Only include countries that have won at least 1 gold in both summer and winter.

This function should return a single string value.

```
In [5]: def answer_three():
        df_copy = df[(df['Gold']>0) & (df['Gold.1']>0)]
        df_copy['diff_rel'] = (df_copy['Gold'] - df_copy['Gold.1']) / df_copy['Gold.2']
        return df_copy['diff_rel'].idxmax()

        answer_three()
```

```
Out[5]: 'Bulgaria'
```

Question 4

Write a function that creates a Series called "Points" which is a weighted value where each gold medal (Gold.2) counts for 3 points, silver medals (Silver.2) for 2 points, and bronze medals (Bronze.2) for 1 point. The function should return only the column (a Series object) which you created.

This function should return a Series named Points of length 146

```
In [6]: def answer_four():  
        df['Points'] = 3*df['Gold.2'] + 2*df['Silver.2'] + 1*df['Bronze.2']  
        return df.Points  
  
        answer_four()  
        #len(df.Points)
```

```

Out[6]: Afghanistan      2
        Algeria          27
        Argentina       130
        Armenia         16
        Australasia      22
        Australia       923
        Austria         569
        Azerbaijan      43
        Bahamas         24
        Bahrain          1
        Barbados         1
        Belarus         154
        Belgium         276
        Bermuda          1
        Bohemia          5
        Botswana         2
        Brazil          184
        British West Indies 2
        Bulgaria        411
        Burundi          3
        Cameroon        12
        Canada          846
        Chile           24
        China           1120
        Colombia        29
        Costa Rica       7
        Ivory Coast      2
        Croatia         67
        Cuba            420
        Cyprus           2
        ...
        Spain           268
        Sri Lanka        4
        Sudan            2
        Suriname         4
        Sweden          1217
        Switzerland     630
        Syria            6
        Chinese Taipei   32
        Tajikistan       4
        Tanzania         4
        Thailand         44
        Togo             1
        Tonga            2
        Trinidad and Tobago 27
        Tunisia         19
        Turkey           191
        Uganda           14
        Ukraine          220
        United Arab Emirates 3
        United States    5684
        Uruguay          16
        Uzbekistan       38
        Venezuela        18
        Vietnam          4
        Virgin Islands   2
        Yugoslavia       171

```

```

Independent Olympic Participants    4
Zambia                             3
Zimbabwe                           18
Mixed team                         38
Name: Points, dtype: int64

```

Part 2

For the next set of questions, we will be using census data from the [United States Census Bureau](http://www.census.gov/popest/data/counties/totals/2015/CO-EST2015-alldata.html) (<http://www.census.gov/popest/data/counties/totals/2015/CO-EST2015-alldata.html>). Counties are political and geographic subdivisions of states in the United States. This dataset contains population data for counties and states in the US from 2010 to 2015. [See this document](http://www.census.gov/popest/data/counties/totals/2015/files/CO-EST2015-alldata.pdf) (<http://www.census.gov/popest/data/counties/totals/2015/files/CO-EST2015-alldata.pdf>) for a description of the variable names.

The census dataset (census.csv) should be loaded as census_df. Answer questions using this as appropriate.

Question 5

Which state has the most counties in it? (hint: consider the sumlevel key carefully! You'll need this for future questions too...)

This function should return a single string value.

```

In [7]: census_df = pd.read_csv('census.csv')
        census_df.head()

```

```

Out[7]:

```

	SUMLEV	REGION	DIVISION	STATE	COUNTY	STNAME	CTYNAME	CENSUS2010POI
0	40	3	6	1	0	Alabama	Alabama	4779736
1	50	3	6	1	1	Alabama	Autauga County	54571
2	50	3	6	1	3	Alabama	Baldwin County	182265
3	50	3	6	1	5	Alabama	Barbour County	27457
4	50	3	6	1	7	Alabama	Bibb County	22915

5 rows × 100 columns

```

In [8]: def answer_five():
        temp = census_df.groupby('STNAME').agg('sum')
        return temp.COUNTY.idxmax()
        answer_five()

```

```

Out[8]: 'Texas'

```

Question 6

Only looking at the three most populous counties for each state, what are the three most populous states (in order of highest population to lowest population)? Use CENSUS2010POP.

This function should return a list of string values.

```
In [81]: filtered_df = census_df[census_df['SUMLEV']==50]
state_df = pd.DataFrame()
state_df['State'] = filtered_df.STNAME.unique()

state_df['Top3PoP'] = 0
state_df.set_index('State', inplace=True)

for st in state_df.index:
    countiespop = filtered_df[census_df['STNAME'] == st].sort(['CENSUS2010POP'], ascending=False)['CENSUS2010POP']
    # countiespop = interim[:3].sum()
    if type(countiespop) == pd.Series:
        stsum = countiespop.iloc[:3].sum()
    else:
        stsum = countiespop
    state_df['Top3PoP'].loc[st] = stsum
    # if type(countiespop) == pd.Series:
    #     stsum = sum()

state_df = state_df.sort(['Top3PoP'], ascending=False)
```

```
In [82]: def answer_six():
        return state_df[:3].index.tolist()

answer_six()
```

```
Out[82]: ['California', 'Texas', 'Illinois']
```

Question 7

Which county has had the largest absolute change in population within the period 2010-2015? (Hint: population values are stored in columns POPESTIMATE2010 through POPESTIMATE2015, you need to consider all six columns.)

e.g. If County Population in the 5 year period is 100, 120, 80, 105, 100, 130, then its largest change in the period would be $|130-80| = 50$.

This function should return a single string value.


```
In [107]: popcols = ['POPESTIMATE20'+ str(x) for x in range(10,16)]
cols = ['SUMLEV', 'CTYNAME'] + popcols
dfp = census_df[cols]
dfp = dfp[dfp['SUMLEV'] == 50]
dfp.set_index('CTYNAME', inplace=True)
maxx = dfp[popcols].max(axis=1)
minn= dfp[popcols].min(axis=1)
dfp['diff'] = (maxx - minn).abs()
dfp = dfp.sort('diff', ascending=False)
```

```
In [108]: def answer_seven():
            return dfp['diff'].idxmax()

answer_seven()
```

Out[108]: 'Harris County'

Question 8

In this datafile, the United States is broken up into four regions using the "REGION" column.

Create a query that finds the counties that belong to regions 1 or 2, whose name starts with 'Washington', and whose POPESTIMATE2015 was greater than their POPESTIMATE 2014.

This function should return a 5x2 DataFrame with the columns = ['STNAME', 'CTYNAME'] and the same index ID as the census_df (sorted ascending by index).

```
In [11]: def answer_eight():
            temp1 = census_df[(census_df['REGION']==1) |
(census_df['REGION']==2)]
            # temp1[temp1['CTYNAME'].str.startswith('Washington', na=False)]
            result = temp1[(temp1['CTYNAME'].str.startswith('Washington', na=False)) & \
                (temp1['POPESTIMATE2015'] > temp1['POPESTIMATE2014'])].loc[:,
['STNAME', 'CTYNAME']]
            return result

answer_eight()
```

Out[11]:

	STNAME	CTYNAME
896	Iowa	Washington County
1419	Minnesota	Washington County
2345	Pennsylvania	Washington County
2355	Rhode Island	Washington County
3163	Wisconsin	Washington County

In []: